

Title: LambdaNetworks: Modeling long-range Interactions without Attention

<https://openreview.net/pdf?id=xTJEN-ggl1b>

Authors: Under double blind review

Field: Computer Vision

Submitted: Sep 28

Modelling Long-Range Interactions:

Notation: \mathbf{q} denotes a query, \mathbf{c} denotes a context.

x is a scalar, \mathbf{x} is a vector, X is a tensor.

$\dim \mathbf{Q}_n = k$ (key depth, eg. embedding dim)

$\dim \text{value}(\mathbf{Y}_n) = v$ (value depth, eg. embedding dim)

intra depth = $u=1$ for classical transformers.

$|m|$ is the dimensionality of a set indexed by m

$(\mathbf{Q}_n, \mathbf{C}_m)$ denotes a content based relation (function or assignment)

$(\mathbf{Q}_n, (\mathbf{n}, \mathbf{m}))$ denotes a position based relation.

Table 1: Hyperparameter, parameters and quantities of interest describing our lambda layer.

Name	Type	Description
$ k , v , u $	hyperparameter	key/query depth, value depth, intra-depth
$\mathbf{W}_Q \in \mathbb{R}^{d \times k }$ $\mathbf{W}_K \in \mathbb{R}^{d \times k \times u }$ $\mathbf{W}_V \in \mathbb{R}^{d \times v \times u }$ $\mathbf{E}_{nm} \in \mathbb{R}^{ k \times u }$	parameter	a tensor that linearly projects the inputs a tensor that linearly projects the context a tensor that linearly projects the context a positional embedding for the relation (n, m) .
$\mathbf{X} \in \mathbb{R}^{ n \times d}$ $\mathbf{C} \in \mathbb{R}^{ m \times d}$	input	the inputs the context
$\mathbf{Q} = \mathbf{X}\mathbf{W}_Q \in \mathbb{R}^{ m \times k \times u }$ $\mathbf{K} = \mathbf{C}\mathbf{W}_K \in \mathbb{R}^{ m \times k \times u }$ $\mathbf{V} = \mathbf{C}\mathbf{W}_V \in \mathbb{R}^{ m \times v \times u }$ $\bar{\mathbf{K}} = \text{softmax}_m(\mathbf{K})$	activation	the queries the keys the values the normalized keys
$\mu_m^c = \mathbf{K}_m \mathbf{V}_m^T \in \mathbb{R}^{ k \times v }$ $\mu_{nm}^p = \mathbf{E}_{nm} \mathbf{V}_m^T \in \mathbb{R}^{ k \times v }$		<i>content</i> contribution from context element m <i>position</i> contribution from context element m
$\mathbf{Y} \in \mathbb{R}^{ n \times d}$	outputs	the outputs

Keys:

Content based relations requires a k dimensional vector \mathbf{k}_m (key) for each context

Position based relations require a k dimensional embedding vector \mathbf{e}_{mn} for each position (m, n) .

The output $\mathbf{Y} = \{\mathbf{Y}_n\}$ does not have a dimension of key depth v or position index m ; these indices must be contracted (matrix multiply, dot product, diag(matrix) etc) away.

Classical attention contracts key and query over a common depth dimension.

Lambda layer maps each query to an output via a linear function $\mathbf{y}_n = \lambda(\mathbf{C})_n(\mathbf{q}_n)$. Each lambda function is computed once, then acts independently of the context. It is then discarded after being applied to its query.

Lambda layer:

The lambda function has both a context and position part (equivalent to position embedding) It takes the form of a k by v matrix.

$$\begin{aligned}\lambda^c &= \sum_m softmax(K_m) V_m \\ \lambda_n^p &= \sum_m E_{nm} V_m \\ \lambda_n &= \lambda^c + \lambda_n^p\end{aligned}$$

The input is transformed to a query via a weight matrix $\mathbf{q}_n = \mathbf{W}_q \mathbf{x}_n$. Then $\mathbf{y}_n = \lambda_n \mathbf{q}_n$ via a matrix multiplication.

Translation Equivalence:

To implement translation equivalence, define a relative position embedding for each (n,m) $R_r(m, n) = E_{m,n}$ where R is a (k,|r|,u) tensor and E is a (k,|n|,|m|,u) tensor. Where r indexes relative positions only.

Lambda Convolution:

For local attention, can generate a convolution with position lambda λ_p kernel. This can be done for Conv2D given 2d context, and is highly optimized.

Lambda convolution has linear time and space complexity wrt. to input length |m|.

Complexity

Batch size b

Time $O(bknmv)$

Space $O(bknv + kmn)$

Still quadratic (mn term). Note however the space complexity has no bmn term; the quadratic term does not scale with batch size, this allows for the processing of large batches.

Multiquery lambda:

Analogous to multi-head attention. However, an important distinction; multi-head attention increases representation power and complexity, multiquery lambda decreases representation power and complexity.

Procedure: split model dimension d into $h \times v$.

Results:

Table 3: **Comparison of the lambda layer and attention mechanisms on ImageNet classification with a ResNet50 architecture.** The lambda layer strongly outperforms alternatives at a fraction of the parameter cost. We include the reported improvements compared to the ResNet50 baseline in subscript to account for training setups that are not directly comparable. [†]: Our implementation.

Layer	Params (M)	top-1
Conv (He et al., 2016) [†]	25.6	76.9 _{+0.0}
Conv + channel attention (Hu et al., 2018b) [†]	28.1	77.6 _{+0.7}
Conv + double attention (Chen et al., 2018)	33.0	77.0
Conv + efficient attention (Shen et al., 2018)	-	77.3 _{+1.2}
Conv + relative self-attention (Bello et al., 2019)	25.8	77.7 _{+1.3}
Local relative self-attention (Ramachandran et al., 2019)	18.0	77.4 _{+0.5}
Local relative self-attention (Hu et al., 2019)	23.3	77.3 _{+1.0}
Local relative self-attention (Zhao et al., 2020)	20.5	78.2 _{+1.3}
Lambda layer	15.0	78.4 _{+1.5}
Lambda layer ($ u =4$)	16.0	78.9 _{+2.0}

Table 4: **The lambda layer reaches higher accuracies while being faster and more memory-efficient than self-attention alternatives.** Inference throughput is measured on 8 TPUv3 cores for a ResNet50 architecture with input resolution 224x224.

Layer	Complexity	Memory (GB)	Throughput	top-1
Global self-attention	$\Theta(b l h n^2)$	120	OOM	OOM
Axial self-attention	$\Theta(b l h n \sqrt{n})$	4.8	960ex/s	77.5
Local self-attention (7x7)	$\Theta(b l h n m)$	-	440ex/s	77.4
Lambda layer	$\Theta(l k n^2)$	0.96	1160ex/s	78.4
Lambda layer (shared embeddings)	$\Theta(k n^2)$	0.31	1210ex/s	78.0
Lambda layer ($ k =8$)	$\Theta(l k n^2)$	0.48	1640 ex/s	77.9
Lambda convolution (7x7)	$\Theta(l k n m)$	-	1100ex/s	78.1

Table 5: LambdaResNets improve upon the parameter-efficiency of large EfficientNets.

Architecture	Params (M)	top-1
EfficientNet-B6	43	84.0
LambdaResNet152	35	84.0
LambdaResNet200	42	84.3

Table 6: LambdaResNets improve upon the flops-efficiency of large EfficientNets.

Architecture	Flops (G)	top-1
EfficientNet-B6	38	84.0
LambdaResNet-270	34	84.0

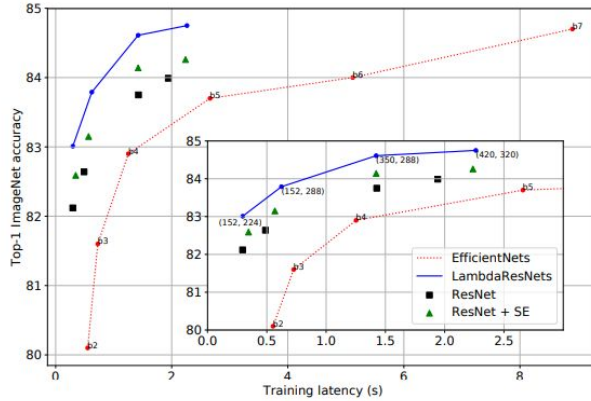


Figure 2: LambdaResNets are $\sim 4.5\times$ faster than EfficientNets and substantially improve the speed-accuracy tradeoff of image classification models³ across different (depth, image size) scales.

Implementation:

$$\begin{aligned}
\lambda_{bkv}^c &= \text{einsum}(\bar{K}_{bmku}, V_{bmvu}) \\
\lambda_{bnkv}^p &= \text{einsum}(E_{knmu}, V_{bmvu}) \\
Y_{bnhv}^c &= \text{einsum}(Q_{bnhk}, \lambda_{bkv}^c) \\
Y_{bnhv}^p &= \text{einsum}(Q_{bnhk}, \lambda_{bnkv}^p) \\
Y_{bnhv} &= Y_{bnhv}^c + Y_{bnhv}^p
\end{aligned}$$

Masked Context Lambda:

Masked attention is achieved by zeroing out certain elements of the attention tensor.

For lambda:

$$\begin{aligned}
\mu_{bmkv}^c &= \text{einsum}(K_{bmku}, V_{bmvu}) \\
\lambda_{bnkv}^c &= \text{einsum}(P_{nm}, \mu_{bmkv}^c) \\
\lambda_{bnkv}^p &= \text{einsum}(E_{knmu} * P_{nm}, V_{bmvu})
\end{aligned}$$

where $p_{nm} = 1[m \in \mathcal{C}_n]$ and $*$ is a broadcasted element-wise multiplication.

p_{mn} is a masking tensor (forward or causal) i.e. is 1 when in context, 0 when out of context(masked).